

OCR for Chinese City Parking Dataset

by Alexander Potekhin
Date: 25/12/2022

1. Setup

1.1. Load data

```
In [ ]: import torch
import torch.nn as nn
import PIL
import numpy as np

from matplotlib import pyplot as plt
from torchmetrics.functional import char_error_rate
from tqdm.notebook import tqdm
from torchvision import transforms, models
from torch.utils.data import Dataset, DataLoader
from PIL import Image
from pathlib import Path
from IPython.display import clear_output

In [ ]: TRAIN_DIR = Path("./CCPD2019-d11/train")
TEST_DIR = Path("./CCPD2019-d11/test")

DEVICE = torch.device("cuda")

In [ ]: train_val_files = list(TRAIN_DIR.rglob('*.jpg'))
test_files = list(TEST_DIR.rglob('*.jpg'))
print("Train files -- ", len(train_val_files))
print("Test files -- ", len(test_files))

Train files -- 199980
Test files -- 9999

In [ ]: # checking all samples have the same length

for file in test_files:
    assert len(file.stem[-7:]) == 7

In [ ]: # getting all possible characters from train files

characters = []
for train_file in train_val_files:
    for train_file_letter in list(train_file.stem[-7:]):
        if train_file_letter not in characters:
            characters.append(train_file_letter)

In [ ]: characters = sorted(characters)
''.join(characters)

Out[ ]: '0123456789ABCDEFHGJKLMNPQRSTUWXYZ云京冀吉宁川新晋桂沪津浙渝湘琼甘皖粤苏蒙藏豫贵陕辽鄂闽陕青鲁黑'

In [ ]: len(characters)

Out[ ]: 66
```

1.2. Visualizing tools

```
In [ ]: %matplotlib inline
plt.style.use('seaborn-v0.8-whitegrid')

In [ ]: def imshow(inp, title=None, plt_ax=plt, preds=None, default=False):
    """Imshow for tensors"""
    inp = inp.numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    inp = std * inp + mean
    inp = np.clip(inp, 0, 1)
    plt_ax.imshow(inp)
    if title:
        plt_ax.set_title("Ground truth: " + title, fontsize=14)
    if preds:
        plt_ax.set_xlabel("Predicted: " + preds, fontsize=14)
    plt_ax.grid(False)

In [ ]: plt.rcParams['font.family'] = ['SimHei', 'sans-serif']
plt.rcParams['axes.unicode_minus'] = False

def draw_samples(test_dataset, preds=None, error_idxs=None):
    """Draw samples from dataset with 3 options:
    1. Samples with ground truth title
    2. Samples with ground truth and predicted title
    3. Samples with ground truth and mispredicted title
    """
    fig, ax = plt.subplots(nrows=3, ncols=3, figsize=(12, 8),
                           sharey=True, sharex=True)
    if error_idxs:
        for fig_ax, error_idx in zip(ax.flatten(), error_idxs):
            pred = preds[error_idx]
            in_val, label = test_dataset[error_idx]
            img_label = label
            imshow(in_val.data.cpu(),
                  title=img_label, plt_ax=fig_ax, preds=pred)
    elif preds:
        for fig_ax in ax.flatten():
            random_characters = int(np.random.uniform(0, 1000))
            pred = preds[random_characters]
            in_val, label = test_dataset[random_characters]
            img_label = label
            imshow(in_val.data.cpu(),
                  title=img_label, plt_ax=fig_ax, preds=pred)
    else:
        for fig_ax in ax.flatten():
            random_characters = int(np.random.uniform(0, 1000))
            in_val, label = test_dataset[random_characters]
            img_label = label
            imshow(in_val.data.cpu(),
                  title=img_label, plt_ax=fig_ax)
```

1.3. Label Encoder

```
In [ ]: class LabelEncoder():
    """Label encoder and decoder for CTC loss
    Insert 'blank' token in the characters for CTC-loss
    """
    def __init__(self, characters):
        self.characters = characters + '-' # for '-1' index
        self.char2idx = {}

        for i, char in enumerate(characters):
            self.char2idx[char] = i + 1 # 0 ind is reserved for 'blank' in CTC-loss

        self.idx2char = {idx: char for char, idx in self.char2idx.items()}

    def encode_sample(self, text):
        """Converts string label to indexes"""
        length = [self.char2idx[char] for char in text]
        text = [len(text)]
        return (torch.IntTensor(text), torch.IntTensor(length))

    def encode(self, batch):
        """Converts labels in batch to indexes"""
        length = [len(sample) for sample in batch]
        text = ''.join(batch)
        text = [self.char2idx[char] for char in text]
        return (torch.IntTensor(text), torch.IntTensor(length))

    def decode_sample(self, text, length):
        """Converts indexes to string labels"""
        chars = []
        for i in range(length):
            if text[i] != 0 and (not (i > 0 and text[i - 1] == text[i])):
                chars.append(self.characters[text[i] - 1])
        return ''.join(chars)

    def decode(self, batch, length):
        """Converts indexes in batch to labels"""
        ind = 0
        texts = []
        for i in range(length.numel()):
            sample_len = length[i]
            texts.append(self.decode_sample(
                batch[ind: ind + sample_len], sample_len))
            ind += sample_len
        return texts
```

2. Dataset

```
In [ ]: class CCPD(Dataset):
    def __init__(self, files):
        super().__init__()
        self.files = sorted(files)

    def __len__(self):
        return len(self.files)

    def load_sample(self, file):
        label = file.stem[-7:] # get label: 1126-皖ASU628.jpg --> 皖ASU628
        img = Image.open(file)
        return img, label

    def __getitem__(self, index):
        transform = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]),
            transforms.Resize(size=(32, 96))
        ])
        x, y = self.load_sample(self.files[index])
        x = transform(x)
        return x, y

In [ ]: #splitting train files to train and val datasets

train_size = int(len(train_val_files) * 0.8)
val_size = len(train_val_files) - train_size

train_dataset, val_dataset = torch.utils.data.random_split(
    CCPD(train_val_files), [train_size, val_size])
test_dataset = CCPD(test_files)

In [ ]: batch_size = 64


train_dataloader = DataLoader(
    train_dataset, batch_size=batch_size, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_size, shuffle=True)
test_dataloader = DataLoader(
    test_dataset, batch_size=batch_size, shuffle=False)

In [ ]: draw_samples(test_dataset)
```

3. Define Model

3.1. Build model

I tried to recreate model from given img



```
In [ ]: class CRNN(nn.Module):
    def __init__(self, num_characters):
        super().__init__()
        self.num_characters = num_characters

        # [batch_size, 3, 32, 96]
        self.conv1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=8, kernel_size=3, padding=2),
            nn.BatchNorm2d(8),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2)
        )
        # [batch_size, 8, 17, 49]
        self.conv2 = nn.Sequential(
            nn.Conv2d(in_channels=8, out_channels=16,
                      kernel_size=3, padding=2),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2)
        )
        # [batch_size, 16, 9, 25]
        self.conv3 = nn.Sequential(
            nn.Conv2d(in_channels=16, out_channels=32,
                      kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2)
        )
        # [batch_size, 32, 4, 12]
        self.lstm = nn.Sequential(
            nn.LSTM(32 * 4, 512, batch_first=True, num_layers=2,
                    dropout=0.2, bidirectional=True),
        )
        # [batch_size, 12, 1024]
        self.clf = nn.Linear(1024, num_characters)

        # [12, batch_size, 67]

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = self.conv3(x)
        b, c, h, w = x.size()
        x, _ = self.rnn(x.view(b, w, c * h))
        x = self.clf(x)
        return x.permute(1, 0, 2)

In [ ]: def accuracy(preds, y):
    accuracy = np.mean(np.array(preds) == np.array(y))
    return accuracy

In [ ]: def decode_preds(logits, text, label_encoder):
    """Decode logits into labels"""
    logits = logits.softmax(2).argmax(2).cpu() # --> [12, 64]
    logits = logits.permute(1, 0).numpy() # --> [64, 12]

    masked_tokens = []
    for label_decoded in logits:
        chars = []
        for char_idx in label_decoded:
            chars.append(
                label_encoder.idx2char[char_idx] if char_idx != 0 else '-')
        masked_tokens.append(chars)

    preds = []
    for label_encoded in masked_tokens:
        pred = []
        for idx, char in enumerate(label_encoded):
            if char != label_encoded[idx-1] and char != '-':
                pred.append(char)
        preds.append(''.join(pred))
    return preds

In [ ]: def get_loss(logits, text, label_encoder, criterion):
    """Calculate CTC-loss"""
    input_len, batch_size, vocab_size = logits.size()
    logits = logits.log_softmax(2)
    enc_text, text_lens = label_encoder.encode(text)
    logits_lens = torch.full(
        size=(batch_size,), fill_value=input_len, dtype=torch.int32).to(DEVICE)
    loss = criterion(logits, enc_text, logits_lens, text_lens)
    return loss

In [ ]: def train_loop(model, data_loader, label_encoder, criterion, optimizer, clip_norm):
    model.train()
    losses = []
    for img, text in tqdm(data_loader):
        optimizer.zero_grad()
        logits = model(img.to(DEVICE))
        loss = get_loss(logits, text, label_encoder, criterion)
        loss.backward()
        nn.utils.clip_grad_norm_(model.parameters(), clip_norm)
        optimizer.step()

        preds = decode_preds(logits, label_encoder)
        acc = accuracy(preds, text)

        losses.append(loss.item())

    return losses, acc

In [ ]: def val_loop(model, data_loader, label_encoder, criterion):
    model.eval()
    losses = []
    with torch.inference_mode():
        for img, text in tqdm(data_loader):
            logits = model(img.to(DEVICE))
            loss = get_loss(logits, text, label_encoder, criterion)

            preds = decode_preds(logits, label_encoder)
            acc = accuracy(preds, text)

            losses.append(loss.item())

    return losses, acc

In [ ]: def evaluate(test_dataloader, label_encoder):
    model.eval()
    losses = []
    predictions = []
    texts = []

    with torch.inference_mode():
        for img, text in tqdm(test_dataloader):
            logits = model(img.to(DEVICE))
            preds = decode_preds(logits, label_encoder)

            predictions.extend(preds)
            texts.extend(text)

    acc = accuracy(predictions, texts)
    cer = char_error_rate(predictions, texts)

    print(f'Test accuracy -- {acc}')
    print(f'Test CER -- {cer}')

    return predictions, texts

In [ ]: def train(model, optimizer, criterion, scheduler, epochs, train_dataloader, val_dataloader, 1
    train_losses = []
    val_losses = []
    train_acc_history = []
    val_acc_history = []

    for epoch in range(epochs):
        train_loss, train_acc = train_loop(model, train_dataloader, label_encoder,
                                           criterion, optimizer, clip_norm)
        val_loss, val_acc = val_loop(
            model, val_dataloader, label_encoder, criterion)
        scheduler.step()

        train_losses.append(np.mean(train_loss))
        val_losses.append(np.mean(val_loss))
        train_acc_history.append(train_acc)
        val_acc_history.append(val_acc)

        clear_output(True)
        fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(10, 5))
        ax[0].plot(train_losses, label='train loss')
        ax[0].plot(val_losses, label='val loss')
        ax[0].set_xlabel('Epoch')
        ax[0].set_title('Losses')
        ax[1].plot(train_acc_history, label='train accuracy')
        ax[1].plot(val_acc_history, label='validation accuracy')
        ax[1].set_xlabel('Epoch')
        ax[1].set_title('Accuracy')
        plt.legend()
        plt.show()

        print(f'Epoch {epoch + 1}')
        print(
            f'Losses: train -- {np.mean(train_loss)}, val -- {np.mean(val_loss)}')
        print(f'Accuracy: train -- {train_acc}, val -- {val_acc}')
```

4. Train model

```
In [ ]: model = CRNN(len(characters) + 1).to(DEVICE) # length of chars + 1 for CTC token
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CTCLoss(blank=0, zero_infinity=True)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=3, gamma=0.1)
epochs = 8
label_encoder = LabelEncoder(''.join(characters))
clip_norm = 5

In [ ]: train(model, optimizer, criterion, scheduler, epochs,
             train_dataloader, val_dataloader, label_encoder, clip_norm)
```



Epoch: 8
Losses: train -- 0.0010650644053341238, val -- 0.010093448391897837
Accuracy: train -- 1.0, val -- 1.0

5. Test model

5.1. Metrics

Now we good to test our model

```
In [ ]: preds, texts = evaluate(test_dataloader, label_encoder)

0% | 0/157 [00:00<, 2it/s]
Test accuracy -- 0.943994399439944
Test CER -- 0.01018673274666871

We have nice and solid scores on both metrics

Best Test accuracy -- 0.95 Best Test CER -- 0.009
```

```
In [ ]: draw_samples(test_dataset, preds)

0 Ground truth: 皖AKD701 Ground truth: 皖AJ697 Ground truth: 皖AE8E26
20 Predicted: 皖AKD701 Predicted: 皖AJ697 Predicted: 皖AE8E26

0 Ground truth: 皖AJ5780 Ground truth: 皖AF697V Ground truth: 皖A018LO
20 Predicted: 皖AJ5780 Predicted: 皖AF697V Predicted: 皖A018LO

0 Ground truth: 皖AS424S Ground truth: 皖AUT816 Ground truth: 皖AJ5780
20 Predicted: 皖AS424S Predicted: 皖AUT816 Predicted: 皖AJ5780
```

5.2. Get errors

Let's see where our model has the worst CER score

```
In [ ]: error_inxs = []

for ind, (pred, text) in enumerate(zip(preds, texts)):
    CER = char_error_rate(pred, text).numpy()
    if CER > 0.5: # empirically selected threshold
        error_inxs.append(ind)

In [ ]: draw_samples(test_dataset, preds, error_inxs)
```

```
0 Ground truth: 皖AB1930 Ground truth: 鄂LLD155 Ground truth: 皖ABJ356
20 Predicted: 闽AG183Q Predicted: 粤L015K5 Predicted: 皖AJ3414

0 Ground truth: 皖AF88BS Ground truth: 苏B271UK Ground truth: 皖A2W003
20 Predicted: 皖AC8P52 Predicted: 浙R71LK Predicted: 皖A3K653

0 Ground truth: 皖AXD167 Ground truth: 皖AD0T89 Ground truth: 皖APY862
20 Predicted: 皖AX01336 Predicted: 皖TY7Z26 Predicted: 皖AFV3538
```

6. Conclusion

It is clearly seen that model mispredicts in case of really fuzzy samples even I could not always recognize.

There is snow on the characters that covers the half of hieroglyphs

To conclude, I have built an OCR model on CCPD dataset with CRNN architecture. Various optimizers, shedulers and hyperparameters were tested to achieve an excellent accuracy and CER scores, but not ideal though.

To improve model understanding of low-poly or distorted images I think we need to work on DATA. So I believe it's feature-extraction problem and maybe we can add augmentation (my approach with augmentation gave lower scores than without) and detect problematic pictures for upsampling before passing them into model. Also I think we can generate more images with defects by adding noise to original and train model on them